

# Persistent State Is Not Enough: A Controlled Study of Planning Horizons in Incremental MPLS Traffic Engineering

Uğur Efe Yiğit

## Abstract

Reinforcement learning (RL) is increasingly applied to traffic engineering (TE), but whether TE decisions are sequential enough to need it is rarely measured. We study incremental MPLS-TE in which routing state persists, one label-switched path may move per five-minute interval, moves are costly and followed by a hold-down, and an action mask enforces failures and protected-class capacity—formally a Markov decision process. On a flow-level 18-router backbone, a masked contextual bandit that regresses only the immediate reward beats MaskablePPO on all 5 training roots (paired return difference 24.9;  $t$ -interval over roots [3.4, 46.4]); a PPO configuration tuned on one root does not change this on four fresh roots. Varying the discount within each learner family points to the horizon rather than the algorithm: return falls as  $\gamma$  grows for PPO and, beyond  $\gamma = 0.5$ , for a bootstrapped Q-learner; long-horizon PPO oscillates at hold-down expiry. A clairvoyant planner that re-plans every interval gains little from looking beyond the next one. Neither learner is the best controller: a per-interval min-max-utilization MILP whose two parameters are chosen on validation seeds beats both. When a move’s effect is delayed by one interval, the bandit fails by construction and falls below PPO, while the MILP stays ahead of every learner. A prior internal study’s bandit result reproduces on a different software stack; its PPO result does not. Persistent state alone does not make a planning horizon useful; the value of lookahead should be measured before a sequential learner is chosen.

## 1 Introduction

Traffic engineering (TE) decides which paths carry which traffic so that a backbone delivers its demand without congestion, SLA violations or needless reconfiguration. In MPLS networks the decision is concrete: each traffic trunk rides one label-switched path, and an operator or controller moves trunks between a few precomputed candidates as load changes [10, 11]. Because traffic varies over the day and failures arrive unannounced, reinforcement learning (RL) is an appealing tool, and deep RL has been applied to routing and TE in many forms [3, 14, 50, 54, 56].

RL earns its machinery—value functions, discounting, credit assignment over long horizons—only if today’s decision changes tomorrow’s options in ways that a myopic decision rule would get wrong. Whether that is true of TE is rarely

measured. The strongest recent TE systems argue the opposite for the standard setting, in which routing is re-optimized from scratch every interval and therefore has no memory: Teal reduces its RL objective to a one-step return [53], and DOTE reports that direct optimization beats an RL-based scheme by a wide margin [44]. Operational MPLS-TE is different in exactly the property those arguments rely on. LSP moves are incremental and signalled, controllers limit how many trunks change at once, stability timers prevent a trunk from moving again too soon, and protected traffic may move only where capacity is guaranteed. The current assignment is then part of the state, a move constrains the future, and the decision problem is a genuine Markov decision process.

This paper asks, for such a formulation, *whether sequential RL is needed, and why or why not*. We study a flow-level MPLS backbone (18 routers, 64 directed links, 17 demands with SLAs, four candidate LSPs each) in which a controller may move one demand per five-minute interval, moves cost reconfiguration effort, a three-interval hold-down follows each move, and an action mask enforces failed links, hold-downs and protected-class capacity. We compare MaskablePPO [30, 48] with a masked neural contextual bandit that regresses only the immediate reward of each action, under identical observations, masks, budgets and paired traffic.

The comparison was first run as a governed study by the project that developed the environment, which found the bandit ahead. We treat that study as a claim to be tested. Our contributions are:

- **A reproduction and its lesson.** Retraining all of the closed study’s runs on different hardware with byte-identical training traffic, the bandit’s result reproduces and PPO’s does not: a change of device and library versions, which alters PPO’s sampling stream and arithmetic, moves PPO’s per-root holdout return by up to 41 points and the bandit’s by at most 4 (§6.1).
- **A controlled comparison of learners and optimizers.** On 5 training roots and 140 disjoint test episodes, a per-interval min-max-utilization MILP controller whose two parameters are selected on validation seeds is the best controller tested; the bandit beats PPO on every root and reaches the return of the untuned MILP controller, with different trade-offs (the MILP attains higher network utility, the bandit reroutes less) (§6.2).
- **A mechanism.** Most of PPO’s deficit (62 %) is reconfiguration cost: on 3 of 5 roots it settles into flipping a

demand at every hold-down expiry, a small immediate cost that, we argue, its long-horizon advantage estimates fail to attribute (§7). Horizon and algorithm are separated by varying the discount within each learner family (§6.3). A PPO configuration tuned on one root, which beats the bandit there, falls below it on all 4 fresh roots (§6.4).

- **A measurement of the value of lookahead.** Re-planning every interval with perfect foresight, a controller gains only a few points from looking beyond the next interval. Open-loop rollouts show where the per-decision non-myopic value lies: with traffic held fixed, the myopic choice is the 24-interval choice in 88 % of states (33 % under the true future), so the non-myopic value visible to a no-change continuation is almost entirely anticipation of exogenous change (§6.5).
- **A regime where the myopic learner fails.** Delaying the effect of a move by one interval—a controlled form of actuation latency—removes the bandit’s signal by construction; the bandit collapses and falls below PPO, while the optimization baseline remains ahead of both (§6.6).

We make no algorithmic or theoretical claim. All code, configurations, per-episode results and the scripts that produce every number, table and figure in this paper are part of the artifact.

## 2 Background and Related Work

**MPLS traffic engineering.** MPLS-TE steers traffic trunks onto explicitly routed label-switched paths (LSPs) [10, 12]; RSVP-TE signals them and moves them with make-before-break [11], and the same “choose one of a few candidate paths” primitive applies to segment-routing policies [21]. Classical TE solves an optimization per traffic matrix—weight setting [22], oblivious and semi-oblivious routing [8, 34], robust TE [51], multicommodity-flow formulations minimizing the maximum link utilization [52]—or adapts online across a fixed set of paths [19, 33]. Production SDN WANs re-solve periodically [28, 31]. TeXCP [33] made the point that online TE must be stable as well as responsive, which is why our formulation prices moves and imposes a hold-down.

**Learning for routing and TE.** RL for routing dates to Q-routing [17]. Deep RL has been applied to intradomain routing [49, 50], utility-maximizing TE [54], critical-flow rerouting [43, 56], GNN-based and multi-agent TE [3, 13, 14, 25, 26, 29]. The strongest recent systems, however, are not sequential: Teal argues that a TE allocation for one interval does not affect future intervals and reduces the RL return to one step [53]; DOTE trains a network end to end on the realized utilization and reports that an RL-based scheme is markedly worse [44]; FIGRET and HARP are one-shot learned TE [4, 38]. Several DRL-TE formulations are, in our reading, contextual bandits in disguise: the state is the exogenous traffic history and actions do not influence it [50, 56]. Our setting differs in the

property those arguments rely on: the routing configuration persists, changes are incremental (one LSP per interval), costly and subject to hold-down, so the current assignment is endogenous state.

**When simple methods win.** Deep-RL results are sensitive to implementation and tuning and need careful statistics [2, 7, 20, 27, 42]; random search matches model-free deep RL on standard control tasks [39]; in networking, a year-long randomized trial found that learned bitrate control did not beat a simple buffer-based rule [55], and a re-examination of RL for optical resource allocation found that improved heuristics match or beat the published agents [18]. Our contribution is not another instance of this pattern but a controlled attribution of it to horizon versus algorithm.

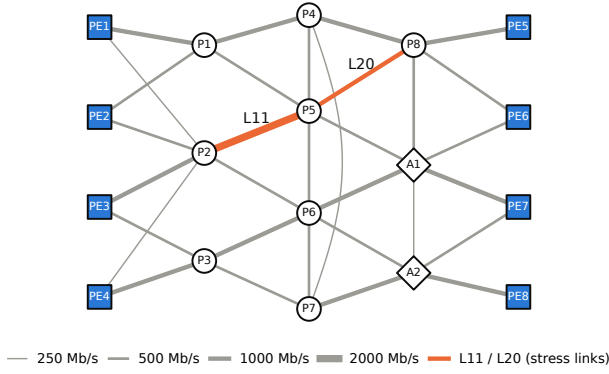
**Contextual bandits and horizon.** A regression oracle on immediate rewards with simple exploration is a standard contextual-bandit design [23, 37], and greedy methods are competitive when contexts are diverse [15]. Theory shows that long horizons need not make RL much harder than bandits in sample complexity [57], while a planning horizon or discount shorter than the true one can improve learned policies with limited data [6, 32, 45]. Laidlaw et al. define the *effective horizon*, the lookahead depth needed to find the optimal action when leaves are evaluated by a base policy, and show it predicts deep-RL success [35, 36]. Congested bandits formalize rewards that depend on recent actions over a short window [9]. Potential-based shaping [41] is policy-invariant only for the discount it is defined with—a caveat for any  $\gamma = 0$  learner trained on a shaped reward. To our knowledge no TE study measures the sequentiality of its decision problem or separates horizon from algorithm in a learner comparison.

**Action masking.** Masking invalid actions in policy-gradient methods is theoretically justified and practically important when many actions are illegal [30]. Masks are a form of safe exploration with external knowledge [24]; constrained MDPs [1, 5] are the alternative we do not pursue. Flow-level simulation may not transfer to packet dynamics [16]; we return to this in Section 9.


## 3 Problem Formulation

We state the decision problem exactly as implemented (Environment V2 of the artifact, frozen at commit dca533b); Appendix A gives every constant.

**Network and demands.** The backbone is a directed graph  $G = (V, E)$  with  $|V| = 18$  routers (four ingress and four egress provider-edge routers, eight core and two aggregation routers) and  $|E| = 64$  directed links (32 bidirectional links, 250 Mb/s to 2,000 Mb/s per direction; Fig. 1). There are  $|\mathcal{D}| = 17$  demands in six classes, each with a priority  $\pi_d \in \{1, \dots, 6\}$ , a delay and a loss SLA, and a *protected* flag (voice and critical traffic). Each demand has a fixed, ordered set of  $k = 4$  candidate



**Figure 1:** Backbone topology: 4 ingress and 4 egress PE routers (squares), 8 core (circles) and 2 aggregation routers (diamonds); line width encodes per-direction capacity. L11 and L20 are the links whose failure or saturation the scripted scenarios stress.



**Figure 2:** Decision timing. Each 5-minute interval is simulated as five one-minute ticks. In the base environment ( $L = 0$ ; the frozen V2 definition) a legal move takes effect at the boundary at which it is chosen, so the immediate reward contains its first interval of effect. The delayed variant ( $L = 1$ ) charges the move at  $t$  and applies it at  $t + 1$ .

label-switched paths  $\mathcal{P}_d = (p_{d,0}, \dots, p_{d,3})$ , computed once by Yen’s algorithm under role, hop-count and propagation-delay constraints;  $p_{d,0}$  is the administratively shortest valid path.

**Exogenous traffic.** The offered rate of demand  $d$  is  $v_d(t) = b_d f_{\kappa(d)}(h(t)) m \eta_d(t) \text{clip}(1 + n_d(t), 0.4, 2)$ : a class-specific diurnal profile, a scenario multiplier, scripted or randomized burst / flash-crowd factors and per-demand AR(1) noise ( $n_d(t) = 0.9 n_d(t-1) + \sigma_d \epsilon$ , one step per simulated minute). Link failures and recoveries are scripted or randomized events. Nothing in the exogenous process  $\xi_t$  depends on routing, which makes comparisons between controllers exactly paired.

**State and transitions.** The controllable state  $y_t$  holds, for each demand, its current candidate  $x_d$ , a disconnection flag, the remaining *dwell* (hold-down)  $\omega_d$ , the previous traffic-engineering (TE) path and its age. The Markov state is  $s_t = (y_t, \xi_t)$ . Decisions are taken every 5 min; an interval is simulated as five one-minute ticks in which a damped fixed-point solver computes carried flow, per-link loss (0 below 90 % utilization, quadratic to 2 % at 100 %,  $1 - 0.98/\rho$  above), M/M/1-shaped queueing delay and per-demand SLA status. A link failure triggers immediate, free fast-reroute of the af-

fected demands to their cheapest live candidate. Because  $\xi$  is action-independent,

$$P(s_{t+1} | s_t, a_t) = P_\xi(\xi_{t+1} | \xi_t) \mathbf{1}[y_{t+1} = \Psi(y_t, a_t, \xi_{t:t+1})], \quad (1)$$

so an action influences the future *only* through the routing state, which persists until changed.

**Actions and masks.**  $\mathcal{A} = \{0\} \cup \{(d, p)\}$ ,  $|\mathcal{A}| = 69$ : action 0 changes nothing;  $(d, p)$  moves demand  $d$  to candidate  $p$ . At most one TE change is made per interval. The legality mask is

$$m_t(d, p) = \mathbf{1}[p_{d,p} \text{ up}] \mathbf{1}[p \neq x_d] \mathbf{1}[\omega_d = 0] \times \mathbf{1}\left[-\text{prot}_d \vee \max_{e \in p_{d,p}} \hat{G}_e / c_e \leq 1\right], \quad (2)$$

with  $m_t(0) = 1$ , where  $\hat{G}_e$  is the projected gross load on link  $e$  after the move. An accepted move starts a three-interval dwell. The last factor couples demands: where one demand is routed changes which moves are legal for protected demands.

**Reward.** With DR the delivered ratio, PD/UD the priority-weighted protected and unprotected disconnected fractions, SEV a bounded priority-weighted SLA severity,  $\rho^{\max}$  the worst link utilization and OV the overload ratio,

$$U_t = 2\text{DR} - 30\text{PD} - 8\text{UD} - 6\text{SEV} - 2g(\rho^{\max}) - 6\text{OV}, \\ r_t = U_t + 0.2(0.995\Phi_{t+1} - \Phi_t) - C_t^{\text{TE}} - 0.05\mathbf{1}[\text{rejected}], \quad (3)$$

where  $g(\rho) = \log(1 + \max(0, (\rho - 0.7)/0.3))$ ,  $\Phi = \tanh(U/10)$  evaluated at interval boundaries, and the move cost  $C^{\text{TE}} = 0.08 + 0.30\text{vol} + 0.12\text{div} + 0.30\text{rev}$  charges the moved traffic share, the edge-set divergence between old and new path and returns to the previous path within six intervals. The reward is the exact sum of twelve logged components.

**What kind of problem is this?** With respect to  $s_t$  the problem is a finite-horizon MDP with state-dependent action sets: routes persist, moves lock a demand for three intervals, reversal costs depend on history, and the protected-class mask couples demands. The learners observe  $o_t \in [0, 1]^{604}$ , which contains the routing state (current and previous path, dwell, path age), per-link utilization, per-demand health relative to SLA, and the projected bottleneck of every candidate move, but not the clock, the noise state or future events; with respect to  $\xi$  the problem is therefore partially observable. A contextual bandit that regresses  $\mathbb{E}[r_t | o_t, a_t]$  discards the persistence of a new configuration beyond the first interval, the opportunity cost of the dwell lock, future reversal costs and the effect on other demands’ legal sets. It keeps the first interval of a move’s effect because, in the base environment, a move carries traffic in the interval in which it is made (Fig. 2). Section 6.5 measures how much lookahead is worth to a controller that replans every interval, and, with a no-change continuation, how much the persistence and cost terms change the ranking of moves; the dwell and coupling terms are probed only through the closed-loop measurement.

**Algorithm 1** Masked value learner (bandit for  $\gamma = 0$ ). Source: `experiments/masked_bandit.py`, `study/learners.py`.

---

```

1: observe  $o$  from 16 environments
2: for each vector step do
3:    $m \leftarrow \text{masks}; a \leftarrow \arg \max_{a:m(a)=1} Q_\theta(o, a)$ , or a uniform legal action w.p.  $\epsilon$ 
4:    $o', r \leftarrow \text{step}(a)$ ; store  $(o, a, m, r, o', m')$ 
5:   if every 4th vector step and  $\geq 4096$  stored then
6:     sample 512 tuples;  $y \leftarrow r$  if  $\gamma = 0$  else Eq. (4)
7:     Adam step on Huber( $Q_\theta(o, a) - y$ ) (selected heads only)
8:    $o \leftarrow o'$ 

```

---

## 4 Methods

All learners see the same observation and the same authoritative mask, train on the same randomized scenario for the same number of environment transitions, and are evaluated deterministically (masked argmax).

**MaskablePPO.** We use the `sb3-contrib` implementation [46] with invalid-action masking [30]: logits of illegal actions are set to  $-10^8$ , so  $\pi_\theta(a | o, m) \propto m(a)e^{z_\theta(o)_a}$ , and the same mask is applied in rollouts, gradient steps and inference. Advantages use GAE [47] and the clipped surrogate objective [48] with value and entropy terms. The study’s default configuration (inherited unchanged from the closed study) is  $\gamma = 0.995$ ,  $\lambda = 0.95$ , learning rate  $3 \cdot 10^{-4}$ , 16 environments  $\times$  512 steps per rollout, batch 512, 8 epochs, clip 0.2, entropy 0.01, separate [256, 256] tanh networks.

**Masked neural contextual bandit.** A 604–256–256–69 ReLU network  $\hat{r}_\theta(o, \cdot)$  predicts the immediate reward of every action. Actions are chosen by masked  $\epsilon$ -greedy ( $\epsilon$  linear from 0.2 to 0.02 over  $2 \cdot 10^5$  transitions; random actions are drawn from the legal set). A replay buffer of the last  $10^5$  tuples  $(o, a, m, r)$  is sampled every 64 transitions for one Adam step on the Huber loss between the *selected* head and the observed reward. There is no next state, discount, target network or bootstrapping. At 400k transitions the bandit performs 6,187 gradient steps, PPO 6,144.

**Masked Q-learner with discount  $\gamma$ .** To vary the planning horizon *without changing anything else*, we extend the bandit by a single parameter. For  $\gamma > 0$  the regression target becomes the normalized double-DQN [40] target

$$y = (1 - \gamma)r + \gamma Q_{\bar{\theta}}(o', \arg \max_{a': m'(a')=1} Q_\theta(o', a')), \quad (4)$$

with a target network copied every 250 updates; the maximization is over legal successor actions only. The factor  $(1 - \gamma)$  rescales all action values of one  $\gamma$  by a common constant and leaves the greedy policy unchanged; it keeps the Huber transition point comparable across  $\gamma$ . With  $\gamma = 0$  the learner is the bandit, bit for bit (unit-tested).

**Non-learning controllers.** *Static* keeps every demand on  $p_{d,0}$ . *Greedy* moves the largest demand on the hottest link (if above 85 %) to the candidate with the lowest bottleneck utilization (margin 5 %). *CSPF* re-replaces demands every six intervals in priority order onto the cheapest candidate with 10 % reservation headroom (hysteresis 8 %). We add *MILP-track*: every interval it solves the classical min-max-utilization path-selection MILP over the four candidates per demand (HiGHS via SciPy), with a small penalty for changing a demand’s path, and, if the optimum lowers the maximum utilization by at least 0.02, submits the single legal move towards it with the largest volume. Because V2 allows one change per interval, only the first legal proposal of any controller is applied.

**Clairvoyant lookahead references.** *Oracle-H* clones the simulator, including the traffic generator, and scores each legal first action by the  $H$ -interval return of taking it and then making no further change; it plays the best (ties: no-op). *Oracle-1* maximizes the exact immediate reward, the quantity the bandit estimates; larger  $H$  adds clairvoyant lookahead. These are references that see the future, not deployable controllers and not optimal policies.

**Controlled delayed effect.** To create a regime in which sequential credit assignment is unavoidable, we add a variant in which a move requested at interval  $t$  is charged at  $t$  but takes effect at  $t + L$  (actuation latency, or a periodic commit cycle). Pending requests lock their demand and are visible in the observation (689 features); a request that has become illegal by  $t + L$  is cancelled. With  $L = 0$  the variant is bit-identical to the base environment (unit-tested). With  $L \geq 1$  the immediate reward of a move contains its cost but none of its effect, so a bandit cannot credit it.

## 5 Experimental Setup

**Training.** Every learner trains on `random_day` (24 h, 288 decisions, randomized bursts, failures and flash crowds per episode) for 400,000 transitions from 16 vectorized environments, with checkpoints every 50,000 transitions. The two learners of a training root share the training episode-seed ledger byte for byte, so roots are paired.

**Evaluation protocol.** Seven scripted scenarios (Table 10) are evaluated with three disjoint seed sets that are fixed in code before any new result was produced: *validation* seeds (2001–2005, 35 episodes) select a checkpoint (highest mean return; ties to the earlier); *test* seeds (3001–3020, 140 episodes) are used only for reporting, for the selected and for the final checkpoint; and *diagnostic* seeds (4001–4003) are used only for structural diagnostics. The closed study’s selection seeds (101–105) and final-holdout seeds (1001–1005) are used only to reproduce its numbers (Section 6.1). Evaluation is deterministic (masked argmax). Because traffic is exogenous, every policy faces byte-identical traffic on a given (scenario, seed), so all comparisons are paired.

**Metric.** The primary metric is the undiscounted operational return  $R = \sum_t r_t$  of an episode, averaged over the 7 scenarios  $\times$  20 seeds. Scenarios have 48–288 decisions, so the average weights long scenarios more; we therefore also report per-scenario results and the operational components (delivered ratio, SLA-violating demand-intervals, maximum utilization, reroutes per hour, moved bandwidth).

**Statistics.** We separate two sources of uncertainty. For a *fixed* policy, 95 % intervals are scenario-stratified bootstrap intervals over test episodes (10,000 resamples). For a *learner*, the experimental unit is the training root: intervals are two-stage bootstrap intervals (resample roots, then scenario-stratified episodes within roots), and we also report a  $t$ -interval over root-level means and the number of roots on which a difference has a given sign. Paired differences are reported in reward units together with the paired standardized effect  $d_z$ . We do not use significance tests as decision rules [2, 42]. The two-stage bootstrap conditions on the roots that were drawn: with two or three roots it stays close to the range of the root means and therefore understates between-root uncertainty. For the five-root primary comparison we report both intervals and lead with the  $t$ -interval; for ablations with two or three roots (horizon sweep, delay, tuning) we report per-root values and sign counts, and the bootstrap interval only as a description of episode-level variability. Per-scenario intervals resample roots, then episodes within roots.

**Hardware.** All new results were produced on a 4-core Intel Xeon (2.1 GHz) virtual machine without GPU, one thread per run (PyTorch 2.14 CPU, SB3/SB3-contrib 2.9.0). The closed study trained on an RTX 4070 laptop GPU; its checkpoints were not archived, so we retrained them (Section 6.1).

## 6 Results

### 6.1 Reproducing the closed study

The study we build on reported, on its untouched final holdout (7 scenarios  $\times$  seeds 1001–1005, three training roots), a mean return of 18.2 for the bandit and 9.0 for PPO, and a small PPO advantage in the deceptive-local-optimum scenario. Its checkpoints were not archived, so we retrained all six runs with the governed trainer at the same commit lineage on a CPU-only machine (the originals used CUDA) and repeated its selection and holdout procedure exactly.

Three properties reproduce *exactly*: the 21 baseline (policy, scenario) holdout means and standard deviations (to  $3 \cdot 10^{-14}$ ), the training episode-seed ledger of every root (byte-identical), and all integrity counters. The bandit reproduces statistically (18.0 vs. 18.2; Table 1); PPO does not (−12.0 vs. 9.0). With identical training traffic, the executions differ in device (CUDA vs. CPU), in library versions (PyTorch, NumPy, Python patch level) and hence in floating-point arithmetic and in PPO’s action-sampling random stream, which SB3 draws from the

**Table 1:** Closed-study reproduction under its own protocol (selection on seeds 101–105, holdout 1001–1005, 35 episodes). Selected checkpoint and mean holdout return, historical (CUDA) vs. reproduced (CPU).

| Learner | Root   | Selected (hist./repr.) | Hist. | Repr. |
|---------|--------|------------------------|-------|-------|
| Bandit  | 42     | 250k / 400k            | 16.1  | 14.7  |
| Bandit  | 271828 | 400k / 350k            | 18.0  | 22.2  |
| Bandit  | 314159 | 300k / 250k            | 20.5  | 17.2  |
| PPO     | 42     | 250k / 350k            | 6.6   | −20.1 |
| PPO     | 271828 | 150k / 400k            | 12.2  | −29.0 |
| PPO     | 314159 | 350k / 200k            | 8.3   | 13.0  |

device’s generator; the bandit’s exploration and replay sampling use a NumPy generator and are device-independent. These differences moved PPO’s per-root holdout return by up to 41 points and the bandit’s by at most 4. The comparison is therefore not a pure numerical-perturbation experiment for PPO; it shows that PPO’s outcome depends strongly on incidental properties of an execution, which the spread across our five CPU training roots confirms (root standard deviation 17.5 for PPO vs. 4.2 for the bandit). On 2 of 3 roots the reproduced PPO never reached a positive validation return at any checkpoint. The bandit’s advantage reappeared on all three roots and grew to 30.1 (per root +34.8 / +51.2 / +4.2, roots 42 / 271828 / 314159; with three roots the  $t$ -interval [−29.2, 89.3] is uninformative about magnitude); the historical deceptive-scenario PPO advantage (−1.1 for the bandit) did not reproduce (bandit − PPO = 14.7, positive on all three roots and in every paired episode).

We draw two methodological conclusions before any new experiment. First, PPO’s outcome in this environment has a large run-to-run variance that is not captured by three roots, so the remainder of the paper uses five roots for the primary comparison and reports root-level uncertainty. Second, the bandit’s result is robust to the perturbation; this is the first indication that its advantage is not an accident of one execution.

### 6.2 RQ1: Does sequential PPO outperform contextual learning?

Table 2 reports all policies on the 140 test episodes. The bandit reaches a mean return of 25.2 [21.5, 28.1] over 5 training roots and PPO 0.3 [−13.7, 13.6] over 5. On paired episodes the bandit’s advantage is 24.9 ( $t$ -interval over roots [3.4, 46.4]; two-stage bootstrap [12.2, 39.1]); it is positive on 5 of 5 roots and in 86 % of paired episodes (paired  $d_z = 0.75$ , pooled over roots and episodes). The ordering does not depend on checkpoint selection: with the final 400k checkpoints the gap is 35.7 (two-stage bootstrap [17.2, 54.5]) (5 of 5 roots). Per scenario (Table 3), the bandit is ahead in every scenario except the double-failure scenario, where the difference is small and its interval includes zero (0.6); the largest gap is on the 24-hour `full_day` scenario. Both learners dominate the static, random and CSPF baselines. Against greedy the bandit is ahead on 5 of 5 roots (21.2 [17.6, 24.1]), PPO on 3 of 5 (−3.7 [−17.6,


9.8]).

We also compare with a controller that does not learn at all. MILP-track, which solves a min-max-utilization path-selection problem on current telemetry every interval and moves one demand towards the solution, reaches 24.0 [21.9, 26.3]. We detect no difference in return between the bandit and MILP-track at five roots (bandit – MILP-track = 1.2,  $t$ -interval  $[-4.0, 6.3]$ , positive on 4 of 5 roots); this is not an equivalence claim, and the two differ by scenario (Table 15): on the double-failure scenario the bandit’s difference is  $-16.6$  and positive on 0 of 5 roots. PPO falls short of MILP-track ( $-23.8$ ,  $t$ -interval  $[-45.5, -2.0]$ , 0 of 5 roots positive). Operationally, the bandit delivers 95.31 % of offered traffic with 164 SLA-violating demand-intervals per episode and 2.30 reroutes per hour, against 94.79 %, 190 and 6.38 for PPO and 95.51 %, 162 and 5.83 for MILP-track; the return gap to PPO thus corresponds to about half a percentage point of delivered traffic, fewer SLA violations and a third of PPO’s churn. MILP-track attains higher network utility and pays more in move costs, which its objective does not contain (Table 7); the bandit reaches a similar return with 0.39 times its reroute rate. MILP-track’s two parameters (the minimum utilization gain that triggers a move, and which mismatched demand to move) were not tuned for this reward. Selecting them on the validation seeds (eight combinations; selected: minimum gain 0.10, the largest value in the grid, and largest-volume demand first; validation return rose with the threshold, so a wider grid might do better still) gives a controller that reaches 30.9 [28.8, 33.0] on the test episodes with 3.77 reroutes per hour, ahead of the bandit by 5.7 ( $t$ -interval  $[0.5, 10.9]$ ; the bandit is ahead on 0 of 5 roots) and of PPO by more. With a model of the objective and a little tuning, per-interval optimization beats both learners of the main comparison here. The comparison is asymmetric—MILP-track was tuned, the learners were not—and one change to the bandit (training without the shaping term, §6.4) changes that: shaping-free bandit – tuned MILP-track =  $-4.2 / +0.3$  on roots 42 / 314159.

### 6.3 RQ2: Is it the horizon or the algorithm?

PPO differs from the bandit in two ways at once: it optimizes a discounted return with  $\gamma = 0.995$ , and it is an on-policy policy-gradient method rather than a value regressor on replayed data. We vary the discount within each family on the three preregistered roots, holding everything else fixed (Table 4, Fig. 3).

*Within PPO*, removing the horizon closes the gap. PPO with  $\gamma = 0$  reaches 23.6 against the bandit’s 24.2 on the same roots (mean difference  $-0.6$ , above the bandit on 2 of 3 roots), while PPO with  $\gamma = 0.995$  reaches  $-6.4$  (above the bandit on 0 of 3 roots); with the intermediate  $\gamma = 0.9$  it reaches 15.4 (above the bandit on 0 of 3 roots), so PPO’s mean return falls as its horizon grows (on two of the three roots individually), and at  $\gamma = 0.9$  it is close to the Q-learner with the same discount (19.4). The oscillation disappears with the horizon: reversals fall from 52.8 to 1.0 per episode, and the validation curves become monotone (Fig. 3). This points to the long horizon,



**Figure 3:** Mean validation return (seeds 2001–2005) during training. Thin lines: individual roots. PPO with  $\gamma = 0.995$  is unstable across roots; with  $\gamma = 0$  it learns as fast as the bandit.

rather than the policy-gradient algorithm, as what limits PPO here. It does not exclude an interaction between PPO and long horizons—for example the scale of undiscounted-sum value targets—which the reward-normalization and GAE configurations of §6.4 probe.

*Within the value family*, adding a horizon does not help. The masked Q-learner is the bandit with one change—a bootstrapped, discounted target (Eq. 4). With  $\gamma = 0.5$  it is close to the bandit (mean difference 1.4, better on 1 of 3 roots); with  $\gamma = 0.9$  it is worse (19.4 vs 24.2; mean difference  $-4.8$ , better than the bandit on 0 of 3 roots), mainly through lower network utility rather than higher churn (3.7 reversals per episode); with  $\gamma = 0.99$  it falls to 3.3 (mean difference  $-20.9$ , better on 0 of 3 roots). In both families the return decreases as the discount grows beyond 0.5. This matches the sequentiality diagnostic: the part of the future that a learner could exploit from its observation is small, so a bootstrapped target adds estimation error without adding much signal—the effect of a shorter planning horizon under limited data [6, 32].

### 6.4 RQ3: Is the gap an artifact of tuning, budget, shaping or the observation?

**PPO tuning (E3).** Eight one-factor variations of PPO’s configuration were trained on root 42 and compared with the default on validation seeds only (Appendix B). Of the 8 configurations, the best by validation return is ent003 (15.3, against  $-19.1$  for the default and 12.2 for the bandit on the same root); Table 12 lists all of them with their test returns, which played no role in selection. On root 42 the selected configuration (entropy coefficient 0.03) is ahead of the bandit on validation and test, so H4 fails on the selection root. Root 42 is also the root on which default PPO was worst, and one-root selection is noisy, so the configuration was retrained on the four other primary roots (E3b). There, tuned PPO minus the bandit is  $-28.6$  (positive on 0 of 4 roots), and tuned minus default PPO is  $-5.5$  (positive on 2 of 4; Table 5). The improvement on

**Table 2:** Test results (7 scenarios  $\times$  20 seeds, 140 paired episodes). Learners: mean over training roots of the selected checkpoint, two-stage bootstrap 95 % CI; fixed policies: scenario-stratified bootstrap CI. Per-interval = return divided by the number of decisions, averaged over episodes. SLA = SLA-violating demand-intervals per episode. † clairvoyant reference (sees future traffic), not a controller.

| Policy                | Return [95 % CI]        | Per-interval | Delivered | SLA | Mean max-util | Reroutes/h | Moved Mb/s |
|-----------------------|-------------------------|--------------|-----------|-----|---------------|------------|------------|
| Masked bandit         | 25.2 [21.5, 28.1]       | −0.302       | 0.9531    | 164 | 1.033         | 2.30       | 1794       |
| MaskablePPO           | 0.3 [−13.7, 13.6]       | −0.577       | 0.9479    | 190 | 1.047         | 6.38       | 4988       |
| No-op                 | −95.5 [−98.0, −92.9]    | −1.602       | 0.9058    | 352 | 1.385         | 0.00       | 0          |
| Static SP             | −96.5 [−98.9, −94.0]    | −1.618       | 0.9040    | 346 | 1.396         | 0.37       | 204        |
| Random valid          | −107.5 [−111.3, −103.6] | −1.744       | 0.8944    | 401 | 1.315         | 11.66      | 9908       |
| CSPF                  | −25.6 [−28.0, −23.2]    | −0.896       | 0.9364    | 243 | 1.195         | 0.61       | 636        |
| Greedy                | 4.0 [1.7, 6.3]          | −0.484       | 0.9485    | 187 | 1.066         | 4.60       | 9293       |
| MILP-track            | 24.0 [21.9, 26.3]       | −0.295       | 0.9551    | 162 | 1.002         | 5.83       | 6312       |
| MILP-track (tuned)    | 30.9 [28.8, 33.0]       | −0.215       | 0.9562    | 161 | 0.994         | 3.77       | 4244       |
| Oracle-1 <sup>†</sup> | 46.9 [45.2, 48.7]       | 0.010        | 0.9598    | 132 | 0.962         | 3.36       | 3233       |

**Table 3:** Bandit – PPO paired return difference per scenario (selected checkpoints; two-stage bootstrap 95 % CI over roots, then episodes), number of roots with a positive difference, and share of paired episodes the bandit wins.

| Scenario       | Difference [95 % CI] | Roots + | Episodes won |
|----------------|----------------------|---------|--------------|
| Deceptive      | 9.1 [2.5, 16.5]      | 5/5     | 81 %         |
| Evening peak   | 33.0 [18.0, 47.9]    | 5/5     | 100 %        |
| Flash crowd    | 21.0 [8.7, 33.2]     | 5/5     | 89 %         |
| Full day       | 69.5 [24.5, 124.0]   | 5/5     | 97 %         |
| Link failure   | 18.4 [7.0, 29.7]     | 5/5     | 87 %         |
| Double failure | 0.6 [−11.4, 10.9]    | 3/5     | 57 %         |
| Overload       | 22.8 [11.0, 31.9]    | 5/5     | 91 %         |

**Table 4:** Horizon sweep (test return, selected checkpoints) on up to three training roots (42, 314159, 271828; column Roots). The value learner with  $\gamma = 0$  is the bandit. “vs. bandit”: mean paired difference to the bandit on the same roots, and per root (roots 42 / 271828 / 314159).

| Learner            | $\gamma$ | Roots | Return | vs. bandit                    | Rer./h | Rev. |
|--------------------|----------|-------|--------|-------------------------------|--------|------|
| Q-learner (bandit) | 0        | 3     | 24.2   | –                             | 2.35   | 1.8  |
| Q-learner          | 0.5      | 3     | 25.6   | +1.4 (+8.0 / −0.9 / −2.9)     | 2.99   | 2.7  |
| Q-learner          | 0.9      | 3     | 19.4   | −4.8 (−1.9 / −6.9 / −5.5)     | 2.90   | 3.7  |
| Q-learner          | 0.99     | 3     | 3.3    | −20.9 (−12.0 / −13.0 / −37.7) | 5.36   | 10.5 |
| PPO                | 0        | 3     | 23.6   | −0.6 (+2.5 / +1.2 / −5.4)     | 1.53   | 1.0  |
| PPO                | 0.9      | 3     | 15.4   | −8.7 (−0.5 / −13.5 / −12.1)   | 3.83   | 7.0  |
| PPO                | 0.995    | 3     | −6.4   | −30.6 (−32.0 / −50.8 / −8.9)  | 8.94   | 52.8 |

the selection root did not transfer; even there it depended on checkpoint selection, since the final checkpoint on root 42 oscillates again. With the tuned configuration PPO still oscillates on two of the four fresh roots. The two variations that target the proposed mechanism behave differently on the selection root: GAE  $\lambda = 0.8$ , which shortens the advantage’s credit horizon while keeping  $\gamma = 0.995$ , gives the second-best validation return and a selected checkpoint that never reverses, but its final checkpoint is far worse; reward normalization helps less. Five of the eight selected checkpoints and six of the eight final ones reverse most of their moves. On one root this is consistent with, but no test of, the advantage-noise hypothesis.

**Table 5:** Control experiments: per-root mean test returns of the compared policies on their common roots, and the paired difference (mean; per root, in increasing order of root seed).

| Comparison                                    | Roots | A    | B    | A – B                                 |
|-----------------------------------------------|-------|------|------|---------------------------------------|
| bandit + clock vs. bandit                     | 2     | 24.0 | 22.0 | +1.9 (+1.9 / +1.9)                    |
| Q $\gamma = 0.9$ + clock vs. Q $\gamma = 0.9$ | 2     | 18.7 | 18.3 | +0.4 (+2.3 / −1.5)                    |
| PPO tuned (ent. 0.03) vs. bandit              | 4     | −1.7 | 26.9 | −28.6 (−40.8 / −22.4 / −38.5 / −12.9) |
| PPO tuned vs. PPO default                     | 4     | −1.7 | 3.8  | −5.5 (−30.5 / +0.3 / +12.4 / −4.1)    |
| bandit, no shaping vs. bandit                 | 2     | 28.9 | 22.0 | +6.9 (+8.4 / +5.4)                    |

**Budget (E5).** (pending)

**Shaping (E6).** The bandit is trained on the shaped reward, and potential-based shaping is policy-invariant only for the discount it was designed for; for a myopic learner it adds one step of lookahead on the potential. Training the bandit with the shaping term removed and evaluating it on the primary reward changes its return by 6.9 (2 of 2 roots higher); the shaping term costs the bandit return rather than helping it, so its advantage over PPO does not rest on shaping, and the shaped bandit of the main comparison is not the strongest myopic learner available (we did not train PPO without shaping).

**Time of day (E7).** The open-loop diagnostic attributes the non-myopic value to anticipation of exogenous change, and the observation has no clock. If that is the missing ingredient for a sequential learner, adding the time of day should help a  $\gamma = 0.9$  Q-learner more than the bandit (H9, exploratory). Adding the clock changes the bandit’s return by 1.9 (2 of 2 roots higher) and the Q-learner’s by 0.4 (1 of 2 roots higher). The clock helps both learners a little and the myopic one at least as much, so in this environment it serves as context for the immediate reward rather than as a basis for planning; H9 is not supported.

## 6.5 RQ4: How sequential is the decision problem?


Formally the problem is an MDP (§3); the question is how much a controller gains by looking beyond the current interval. We answer it first in closed loop, then per decision.

**Closed loop: clairvoyant lookahead as a controller.** Oracle- $H$  re-plans every interval with the true future: it rolls the cloned simulator forward  $H$  intervals from every legal action (the action, then no further change) and takes the best. On the first five test seeds (35 episodes), Oracle-1 reaches 48.6, against 26.8 for the bandit and 24.7 for MILP-track. Extending the lookahead to three or six intervals adds only 2.9 [1.5, 4.2] and 1.4 [0.5, 2.5] (Oracle-6 is not better than Oracle-3). Even with perfect foresight, lookahead beyond one interval is worth a few points to this re-planning controller; knowing the next interval exactly is worth 21.7 over the bandit on the full test set. Oracle-1’s lead over the learners combines an exact model of the reward with knowledge of the next interval’s traffic, bursts and failures; we did not separate the two. For this planner, what matters is the next interval, not the ones after it.

**Per decision: open-loop rollouts.** On 495 states visited by the greedy controller (diagnostic seeds 4001–4003, every fourth decision; 390 of them have 24 remaining intervals), we compare the action that is best for the first interval with the action that is best over  $H$  intervals under the same no-change continuation (Fig. 4); this is the effective-horizon idea of Laidlaw et al. [35] with a no-change base policy. *Under the true future*, the first-interval choice is the 24-interval choice in only 33 % of states and captures 52 % of the 24-interval gain; in 59 % of states the 24-interval best move has a negative immediate advantage. These open-loop values overstate what a controller can realise: a move that pays off only after traffic changes can also be made when the change arrives, which is what the re-planning oracle does, and the closed-loop gains above are small. *Holding traffic and link state at their current values* during the rollouts, agreement is 88 %, the myopic choice captures 97 % of the gain and sacrifice moves occur in 8 % of states, in every scenario (Appendix C). With constant exogenous inputs and no further moves, the  $H$ -interval value of a move is close to  $H$  times its per-interval gain minus its one-off cost, so the frozen agreement measures one thing: how often amortizing a move’s cost over a longer stay changes which move is best. It rarely does. The open-loop gap between the two rollouts is anticipation of exogenous change (traffic growth, bursts, failures and recoveries), which our observation does not provide information for (no clock, no forecast; §6.4 tests adding the clock). Neither rollout can detect value that arises from *later* moves—the opportunity cost of a dwell lock, or a move that enables another under the protected-class rule—because the continuation never moves again. The oracle ladder does not bound that value either: it shows what a planner with the same no-change continuation gains, and a planner with a better continuation could gain more. That Oracle-6 is not better than Oracle-3 is consistent with this continuation bias.

## 6.6 RQ5: When does a myopic learner fail?

In the base environment a move carries traffic in the interval in which it is made, so its first-interval effect is in the immediate reward. The delayed-activation variant ( $L = 1$ ) removes that property: a request is charged when it is made and takes effect



**Figure 4:** Open-loop sequentiality diagnostic (greedy-visited states; from 495 states at  $H = 1$  to 390 at  $H = 24$ , since states near an episode’s end lack  $H$  remaining intervals). Left: probability that the action with the best first-interval reward is also the best over  $H$  intervals (“take it, then change nothing”). Right: share of the best  $H$ -interval gain that the myopic action captures. Solid: rollouts carry the episode’s true future traffic and failures (clairvoyant). Dashed: traffic and link state held at their current values.

one interval later; the observation adds the pending request of every demand (689 features), so the state stays Markov. Two side effects are part of the variant: a demand with a pending request cannot be moved again, which lengthens the effective lock from three to four intervals, and a request that has become illegal at activation is cancelled but keeps its charge (Table 6).

The bandit fails by construction, and the ranking of the two learners reverses. On the same two roots the bandit falls from 22.0 without delay to  $-30.3$ , because its regression target now contains the cost of a move but none of its benefit, and it learns to move less. PPO’s mean is nearly unchanged ( $1.6 \rightarrow 1.2$ ), but this average hides opposite per-root shifts, and on root 42 PPO still flips demands at hold-down expiry under delay; it is ahead of the bandit on 2 of 2 roots (mean 31.5). The masked Q-learner with  $\gamma = 0.9$ , whose target includes the discounted value of the next state and hence the delayed effect, partly recovers:  $-5.6$  under delay against 18.3 without, above the bandit on 2 of 2 roots (mean difference 24.7) but below PPO. With  $\gamma = 0.5$  (an effective horizon of about two intervals) it reaches  $-2.0$  (24.6 without delay; above the bandit on 2 of 2 roots). The non-learning controllers degrade because they react to a state that the delayed move no longer matches (MILP-track 24.0  $\rightarrow$  11.3, greedy 4.0  $\rightarrow$   $-5.1$ ), yet MILP-track remains better than every learner. What this experiment establishes is therefore narrow: a learner whose target is the decision-time reward cannot work when effects are delayed, so some form of delayed credit assignment is needed. It does not show that a long-horizon learner is needed—a target that sums the next  $L + 1$  rewards would suffice in principle, although at this budget neither Q-learner recovers the no-delay return—nor that RL beats optimization-based control in this regime.

## 7 Why PPO Falls Behind

**Where the return goes.** Splitting each episode’s return into network utility (delivery, disconnection, SLA, utilization and overload terms), shaping, and move costs (Table 7) localizes PPO’s deficit. Averaged over roots, the bandit earns a utility of 28.0 and pays  $-2.8$  in move and reversal costs; PPO earns

**Table 6:** Delayed TE activation ( $L = 1$ ), mean test return per training root (roots 42 / 314159; “No delay”: the same roots without delay). Non-learning references: the same 140 test episodes. “vs. bandit”: paired difference to the bandit under delay, per root. Rer./h: reroutes per hour under delay.

| Policy                           | No delay     | $L = 1$       | vs. bandit    | Rer./h |
|----------------------------------|--------------|---------------|---------------|--------|
| Masked bandit                    | 18.3 / 25.8  | -25.8 / -34.9 | -             | 0.84   |
| MaskablePPO ( $\gamma = 0.995$ ) | -13.7 / 16.9 | -4.3 / 6.7    | +21.5 / +41.6 | 6.34   |
| Q-learner ( $\gamma = 0.5$ )     | 26.3 / 22.9  | -9.0 / 5.0    | +16.8 / +39.9 | 1.35   |
| Q-learner ( $\gamma = 0.9$ )     | 16.4 / 20.3  | -10.4 / -0.8  | +15.4 / +34.1 | 3.32   |
| CSPF                             | -25.6        | -27.1         | -             | 0.61   |
| Greedy                           | 4.0          | -5.1          | -             | 3.80   |
| MILP-track                       | 24.0         | 11.3          | -             | 5.88   |
| No-op                            | -95.5        | -95.5         | -             | 0.00   |
| Static SP                        | -96.5        | -96.4         | -             | 0.36   |

**Table 7:** Return decomposition on the test episodes (means over episodes and, for learners, over roots). Utility = delivery + disconnection + SLA + utilization + overload terms; Costs = move and reversal terms.

| Policy                | Return | Utility | Costs | Shaping | Moves | Reversals |
|-----------------------|--------|---------|-------|---------|-------|-----------|
| Masked bandit         | 25.2   | 28.0    | -2.8  | 0.01    | 12.4  | 1.8       |
| MaskablePPO           | 0.3    | 18.4    | -18.1 | 0.01    | 46.9  | 32.6      |
| MILP-track            | 24.0   | 34.3    | -10.2 | 0.01    | 39.6  | 8.8       |
| Greedy                | 4.0    | 14.4    | -10.4 | 0.00    | 31.8  | 11.8      |
| CSPF                  | -25.6  | -24.9   | -0.7  | 0.00    | 3.6   | 0.1       |
| Static SP             | -96.5  | -96.1   | -0.3  | -0.01   | 1.8   | 0.0       |
| No-op                 | -95.5  | -95.5   | 0.0   | -0.01   | 0.0   | 0.0       |
| Oracle-1 <sup>†</sup> | 46.9   | 50.6    | -3.7  | 0.01    | 20.0  | 0.4       |

18.4 and pays -18.1. Of the 24.9-point gap, 15.4 (62 %) is move and reversal cost and 9.5 is network utility. PPO’s utility varies widely across roots (7.8 to 32.2, against 20.9 to 31.6 for the bandit); its best root is above the bandit’s best. On 3 of 5 roots more than half of its accepted moves are *reversals* (up to 80 %), and its per-root move costs range from -36.5 to -1.9, against -3.2 to -2.6 for the bandit. The shaping term contributes at most 0.02 per episode to the evaluated return of any policy, but that does not bound its effect on learning: a learner trained on the shaped reward sees it at every step, and removing it changes the bandit’s return by several points (§6.4).

**A degenerate oscillation.** Figure 5 shows one test episode of the deceptive-local-optimum scenario. After moving at nearly every opportunity in the first hour, PPO settles into flipping one demand between two candidates every three intervals—exactly when the hold-down expires—and pays the reversal price each time. The bandit makes a handful of moves at the burst onset and otherwise holds its configuration. The behaviour is not an artefact of checkpoint selection: 2 of 5 final (400k) PPO checkpoints also reverse more than half of their moves.

**A hypothesis: small immediate costs in noisy advantages.** A move that returns a demand to its previous path costs at least 0.38 reward units once (fixed and reversal terms of  $C^{\text{TE}}$ ,

**Table 8:** What the learners’ actions are worth, from exact simulator rollouts on states each learner visits (test seeds 3001–3003, every third step). Top-1: share of states in which the learner’s action is the exact one-interval optimum; regret: exact one-interval reward lost; no-op: share of states in which the learner / the optimum chooses no-op;  $\Delta G_6$ : six-interval value over no-op (no-op continuation) of the learner’s action and of the exact myopic-best action.


|        | Root   | Top-1 | Regret | No-op (learner/opt.) | $\Delta G_6$ learner | $\Delta G_6$ myopic |
|--------|--------|-------|--------|----------------------|----------------------|---------------------|
| Bandit | 42     | 47 %  | 0.27   | 85 / 49 %            | +0.24                | +1.31               |
| Bandit | 271828 | 57 %  | 0.19   | 83 / 61 %            | +0.20                | +0.80               |
| Bandit | 314159 | 60 %  | 0.18   | 85 / 61 %            | +0.27                | +0.82               |
| PPO    | 42     | 9 %   | 0.63   | 12 / 53 %            | -0.08                | +1.37               |
| PPO    | 271828 | 7 %   | 0.67   | 5 / 49 %             | -0.04                | +1.57               |
| PPO    | 314159 | 6 %   | 0.54   | 8 / 67 %             | -0.25                | +0.66               |

Eq. 3), plus its traffic and divergence terms. PPO’s advantage estimates aggregate, through GAE with  $\gamma\lambda \approx 0.945$ , the fluctuations of a reward whose per-interval mean alone ranges from -3.4 to 1.8 across test episodes, driven by bursts, failures and diurnal load. We hypothesise that a one-off cost of this size is small relative to the variance of the advantage target, so that the policy gradient does not attribute it to the step that incurs it. The bandit’s regression target *is* that step’s reward, so the cost enters its estimate without noise from future intervals. We did not measure advantage noise directly; §6.3 tests the explanation indirectly by removing PPO’s horizon, and §8 discusses what speaks against it.

**What the learners’ actions are worth.** On states visited by each learner on test episodes we compute, by cloning the simulator, the exact one-interval reward of every legal action and the six-interval value of the learner’s action, of the exact myopic-best action, and of no-op (Table 8). Neither learner is close to the exact myopic optimum, but they fail in opposite directions. The bandit is conservative: it chooses no-op in 84 % of states although no-op is optimal in only 57 %, and it matches the optimum in 55 % of states (one-interval regret 0.21). PPO over-acts: it chooses no-op in 8 % of states, matches the optimum in 7 % (regret 0.61), and over six intervals its chosen actions are worth *less than doing nothing* on 3 of 3 roots (mean  $\Delta G_6 = -0.12$  against +0.24 for the bandit and +1.20 for the exact myopic choice on PPO’s states). PPO’s frequent moves therefore do not buy delayed value that a myopic criterion would miss.

## 8 Discussion

**What the results mean.** In this formulation the current LSP assignment is genuine state, and a move constrains the future through hold-downs, reversal costs and coupled protected-class legality. Yet a learner that ignores the future reaches the return of an untuned per-interval optimizer (a tuned one is better still), and a learner trained with an effective horizon of about 200 intervals ( $1/(1-\gamma)$ , 17 hours) does worse than both. The closed-loop oracle ladder shows why there is little to gain from looking ahead: even with perfect foresight, a



**Figure 5:** One test episode (deceptive-local-optimum, seed 3001; learners of root 314159). Top: worst link utilization; bottom: accepted moves ( $\times$  = return to the previous path within six intervals). PPO flips one demand every three intervals, the hold-down period; the bandit makes few moves. Oracle-1 sees the next interval’s traffic and is shown for reference.

controller that re-plans every interval gains only a few points from lookahead beyond the next interval. When a move’s effect is felt in the interval in which it is made, the immediate reward ranks moves almost as a longer criterion would; what is not myopic—anticipating bursts, failures and diurnal change—is not provided by an observation without a clock or forecasts, and adding the clock (§6.4) helps the myopic learner at least as much as a sequential one. A sequential learner pays the price of a long-horizon objective (noisy advantages, unstable optimization) without access to the signal that would make it worthwhile.

**Long horizons can hide small immediate costs.** The concrete failure on three of five PPO roots is instructive for RL practice. Reconfiguration costs are small, immediate and certain; congestion rewards are large, delayed and noisy. With  $\gamma = 0.995$  PPO’s advantage estimates are dominated by the latter and fail to attribute the former, so the policy settles into flipping a demand whenever the hold-down allows. We have not measured the advantage noise directly. The evidence for this explanation is that setting  $\gamma = 0$  removes the oscillation and closes the gap; against it, the critic is not poorly fitted in aggregate: in the 10 PPO runs whose final checkpoint oscillates, its explained variance at the last update is 0.83–0.86, against 0.71–0.73 with  $\gamma = 0$ , so the failure would have to lie in the small residual differences between moves rather than in the value scale; and the PPO roots that do not oscillate fall short through lower network utility rather than churn. The discount factor is thus a decisive hyper-parameter here, and the default  $\gamma \approx 0.99$  that most RL-for-networking work inherits is not a neutral choice when costs and benefits live on different

time scales.

**When is sequential RL justified?** Our evidence points to three conditions, each testable before choosing an algorithm: (i) actions whose effects are delayed relative to the decision, so that the immediate reward does not contain them (§6.6); (ii) exogenous dynamics that are predictable from the observation, so that anticipation can be learned; and (iii) decisions that change the future feasible set substantially. The closed-loop oracle ladder of §6.5 measures what one simple clairvoyant lookahead planner gains, and the frozen rollout measures how much cost amortization changes the ranking of moves, both from a simulator without learning. Both score a move as if nothing changed afterwards, so they cannot see (iii): their gains are what such a planner achieves, not an upper bound on the value of lookahead; a rollout with a reactive continuation is the natural next check. We recommend reporting both before choosing an algorithm.

**When should simpler methods be preferred?** When effects are immediate and a model of the objective is available, a per-interval optimizer is the strongest controller we tested: tuned on validation seeds, MILP-track is ahead of the bandit on every root and reroutes less than the untuned version (3.77 vs. 5.83 per hour, against 2.30 for the bandit). A myopic learner, whose observation contains a one-step projection of every move (§9), reaches the return of the untuned optimizer and is well ahead of PPO. Inference cost does not discriminate at a five-minute interval (the bandit needs 0.112 ms per decision, PPO 0.335 ms, MILP-track 31.7 ms). MILP-track’s objective contains neither move costs nor the reward’s SLA

and loss terms; that a hysteresis threshold alone is enough to make it the best controller suggests that a churn-aware optimizer is the baseline any learned TE controller in this setting should be measured against.

**Networking implications.** Teal and DOTE treat TE as a one-shot problem per traffic matrix. Our results support that treatment in a setting with persistent state, incremental changes and stability timers, where one might have expected sequential structure to matter—as long as a change takes effect within the interval in which it is decided, and in a heavily loaded regime where TE mostly decides which traffic to lose (§9). They also show that churn metrics (reroutes, reversals) must be reported with the objective: PPO’s best root reaches a higher network utility than any bandit root (32.2 vs. at most 31.6) but would be unacceptable to an operator because of its oscillation.

**Generalization.** We make no claim beyond this simulator. Larger topologies, elastic traffic, real traffic matrices and control-plane latency all plausibly increase temporal coupling; the diagnostic and the delayed-effect experiment indicate how to test that.

## 9 Limitations

**Simulator realism.** The model is flow-level: delay and loss are analytic functions of utilization, traffic is exogenous and inelastic, the control plane acts instantaneously, fast reroute is free, and telemetry is exact. Learned routing policies trained in fluid models are known not to transfer to packet-level dynamics [16]; our conclusions are about this model. Inelastic traffic in particular removes a natural source of temporal coupling (congestion that suppresses future demand).

**Load regime.** The test scenarios are heavily loaded: the mean per-interval maximum link utilization is 1.00 for MILP-track, 1.03 for the bandit and 1.38 without TE, and even the best controllers lose about 4.5 % of offered traffic. TE here mostly decides where traffic is lost. In a provisioned backbone (utilization well below one), keeping headroom ahead of failures and peaks is a plausible source of sequential value that this study does not test.

**Formulation.** One trunk may move per interval network-wide and each demand has four fixed candidate paths. These choices create the incremental, stateful character of the problem; a controller that re-optimizes all trunks at once faces the one-shot problem studied by Teal and DOTE instead. We did not vary the number of candidates or the per-interval move limit, because the environment definition is frozen.

**Observation.** The observation contains, for every candidate move, the projected gross bottleneck utilization the move would create—a hand-built one-step model of the network.

It makes the immediate consequence of a move nearly observable and so favours a learner that regresses the immediate reward; we did not ablate it, for compute. It omits the clock and any forecast; §6.4 tests adding the clock.

**Scale and data.** One hand-designed 18-router topology, 17 demands, hand-written diurnal profiles and seven scripted test scenarios; no operator traffic matrices or failure traces. No claim of topology generalization is made.

**Learning.** Five training roots for the primary comparison and two to three for ablations; root-level intervals are wide and are reported. A budget of 400k transitions (about 1,390 simulated days), with one longer run per learner. PPO received a bounded one-factor sensitivity study, the bandit none; neither used graph-structured networks, observation normalization or learning-rate schedules. The Q-learner performs one gradient step per 64 transitions (about 6,200 in total) with a target update every 250 steps, and treats the time-limit end of an episode as terminal; that budget may be too small to propagate value over long horizons, so its  $\gamma > 0$  results are a statement about this budget. Learners are trained on randomized days and tested only on scripted scenarios; a shift that affects the two learners differently cannot be excluded. A discount below the true one can help with limited data [6, 32], so “the myopic learner wins at this budget” does not imply that no sequential structure exists—which is why we measure the structure separately.

**Objective.** All rankings use one hand-weighted operational reward; we report its components and operational metrics alongside, but a different operator objective could reorder close methods.

**References.** The clairvoyant oracles see future traffic and only evaluate “one move, then no change”; they bound  $H$ -step-greedy reasoning, not the optimal policy. MILP-track optimizes gross maximum utilization only.

**Provenance.** The closed study’s checkpoints and per-episode holdout data were not archived; its learner results are compared at the aggregate level it committed.

## 10 Conclusion

For an incremental, masked MPLS traffic-engineering problem with persistent routing state, a masked contextual bandit beats MaskablePPO on every training root and reaches the return of an untuned per-interval MILP optimizer; a validation-tuned version of that optimizer is better than both learners of the main comparison. Varying the discount within each learner family points to the planning horizon rather than the learning algorithm: PPO with  $\gamma = 0$  is not detectably different from the bandit at three roots, PPO with  $\gamma = 0.995$  settles into costly oscillations on most roots, and a bootstrapped value learner

does not improve on the bandit. A clairvoyant planner that re-plans every interval but scores each move as if nothing changed afterwards gains little from looking beyond the next interval, and with traffic held fixed the first interval’s reward already ranks moves almost as such a 24-interval criterion does. These are statements about planners with a no-change continuation, not upper bounds on the value of lookahead. Once a move’s effect is delayed, a learner of the decision-time reward fails. A sequential learner should therefore be justified by measured temporal coupling—delayed effects, or exogenous dynamics that the controller can predict—rather than assumed from the fact that the problem is formally an MDP.

## References

- [1] Joshua Achiam, David Held, Aviv Tamar, and Pieter Abbeel. Constrained policy optimization. In *Proceedings of the 34th International Conference on Machine Learning (ICML)*, volume 70 of *PMLR*, pages 22–31, 2017. URL <https://proceedings.mlr.press/v70/achiam17a.html>.
- [2] Rishabh Agarwal, Max Schwarzer, Pablo Samuel Castro, Aaron Courville, and Marc G. Bellemare. Deep reinforcement learning at the edge of the statistical precipice. In *Advances in Neural Information Processing Systems 34 (NeurIPS 2021)*, 2021. URL <https://arxiv.org/abs/2108.13264>. Outstanding Paper Award. arXiv:2108.13264.
- [3] Paul Almasan, José Suárez-Varela, Krzysztof Rusek, Pere Barlet-Ros, and Albert Cabellos-Aparicio. Deep reinforcement learning meets graph neural networks: Exploring a routing optimization use case. *Computer Communications*, 196:184–194, 2022. doi: 10.1016/j.comcom.2022.09.029. arXiv:1910.07421.
- [4] Abd AlRhman AlQiam, Yuanjun Yao, Zhaodong Wang, Satyajeet Singh Ahuja, Ying Zhang, Sanjay G. Rao, Bruno Ribeiro, and Mohit Tawarmalani. Transferable neural WAN TE for changing topologies. In *Proceedings of ACM SIGCOMM 2024*, pages 86–102, 2024. doi: 10.1145/3651890.3672237.
- [5] Eitan Altman. *Constrained Markov Decision Processes*. Chapman & Hall/CRC, 1999. doi: 10.1201/9781315140223.
- [6] Ron Amit, Ron Meir, and Kamil Ciosek. Discount factor as a regularizer in reinforcement learning. In *Proceedings of the 37th International Conference on Machine Learning (ICML)*, volume 119 of *PMLR*, 2020. URL <https://arxiv.org/abs/2007.02040>.
- [7] Marcin Andrychowicz, Anton Raichuk, Piotr Stańczyk, Manu Orsini, Sertan Girgin, Raphaël Marinier, Léonard Hussenot, Matthieu Geist, Olivier Pietquin, Marcin Michalski, Sylvain Gelly, and Olivier Bachem. What matters for on-policy deep actor-critic methods? A large-scale study. In *International Conference on Learning Representations (ICLR)*, 2021. URL <https://openreview.net/forum?id=nIAxjsniDzg>.
- [8] David Applegate and Edith Cohen. Making intra-domain routing robust to changing and uncertain traffic demands. In *Proceedings of ACM SIGCOMM 2003*, pages 313–324, 2003. doi: 10.1145/863955.863991.
- [9] Pranjal Awasthi, Kush Bhatia, Sreenivas Gollapudi, and Kostas Kollias. Congested bandits: Optimal routing via short-term resets. In *Proceedings of the 39th International Conference on Machine Learning (ICML)*, volume 162 of *PMLR*, pages 1078–1100, 2022. URL <https://proceedings.mlr.press/v162/awasthi22a.html>.
- [10] D. Awduche, J. Malcolm, J. Agogbua, M. O’Dell, and J. McManus. Requirements for traffic engineering over MPLS. RFC 2702, September 1999. URL <https://www.rfc-editor.org/rfc/rfc2702>.
- [11] D. Awduche, L. Berger, D. Gan, T. Li, V. Srinivasan, and G. Swallow. RSVP-TE: Extensions to RSVP for LSP tunnels. RFC 3209, December 2001. URL <https://www.rfc-editor.org/rfc/rfc3209>.
- [12] Daniel O. Awduche. MPLS and traffic engineering in IP networks. *IEEE Communications Magazine*, 37(12):42–47, 1999. doi: 10.1109/35.809383.
- [13] Guillermo Bernárdez, José Suárez-Varela, Albert López, Bo Wu, Shihan Xiao, Xiang Cheng, Pere Barlet-Ros, and Albert Cabellos-Aparicio. Is machine learning ready for traffic engineering optimization? In *Proceedings of the 29th IEEE International Conference on Network Protocols (ICNP)*, 2021. doi: 10.1109/ICNP52444.2021.9651930.
- [14] Guillermo Bernárdez, José Suárez-Varela, Albert López, Xiang Shi, Shihan Xiao, Xiang Cheng, Pere Barlet-Ros, and Albert Cabellos-Aparicio. MAGNETO: A graph neural network-based multi-agent system for traffic engineering. *IEEE Transactions on Cognitive Communications and Networking*, 9(2): 494–506, 2023. doi: 10.1109/TCCN.2023.3235719.
- [15] Alberto Bietti, Alekh Agarwal, and John Langford. A contextual bandit bake-off. *Journal of Machine Learning Research*, 22 (133):1–49, 2021. URL <https://jmlr.org/papers/v22/18-863.html>.
- [16] Andreas Boltres, Niklas Freymuth, Patrick Jahnke, Holger Karl, and Gerhard Neumann. Learning sub-second routing optimization in computer networks requires packet-level dynamics. *Transactions on Machine Learning Research*, 2024. URL <https://arxiv.org/abs/2410.10377>. arXiv:2410.10377.
- [17] Justin A. Boyan and Michael L. Littman. Packet routing in dynamically changing networks: A reinforcement learning approach. In *Advances in Neural Information Processing Systems 6 (NIPS 1993)*, pages 671–678, 1993. URL [https://papers.nips.cc/paper\\_files/paper/1993/hash/4ea06fbc83cdd0a06020c35d50e1e89a-Abstract.html](https://papers.nips.cc/paper_files/paper/1993/hash/4ea06fbc83cdd0a06020c35d50e1e89a-Abstract.html).
- [18] Michael Doherty, Robin Matzner, Rasoul Sadeghi, Polina Bayvel, and Alejandra Beghelli. Reinforcement learning for dynamic resource allocation in optical networks: Hype or hope? *Journal of Optical Communications and Networking*, 17(9):D1–D17, 2025. doi: 10.1364/JOCN.559990. arXiv:2502.12804.
- [19] Anwar Elwalid, Cheng Jin, Steven Low, and Indra Widjaja. MATE: MPLS adaptive traffic engineering. In *Proceedings of IEEE INFOCOM 2001*, volume 3, pages 1300–1309, 2001. doi: 10.1109/INFCOM.2001.916625.

- [20] Logan Engstrom, Andrew Ilyas, Shibani Santurkar, Dimitris Tsipras, Firdaus Janoos, Larry Rudolph, and Aleksander Madry. Implementation matters in deep RL: A case study on PPO and TRPO. In *International Conference on Learning Representations (ICLR)*, 2020. URL <https://openreview.net/forum?id=r1etN1rtPB>.
- [21] C. Filsfils, S. Previdi, L. Ginsberg, B. Decraene, S. Litkowski, and R. Shakir. Segment routing architecture. RFC 8402, July 2018. URL <https://www.rfc-editor.org/rfc/rfc8402>.
- [22] Bernard Fortz and Mikkel Thorup. Internet traffic engineering by optimizing OSPF weights. In *Proceedings of IEEE INFOCOM 2000*, volume 2, pages 519–528, 2000. doi: 10.1109/INFCOM.2000.832225.
- [23] Dylan J. Foster and Alexander Rakhlin. Beyond UCB: Optimal and efficient contextual bandits with regression oracles. In *Proceedings of the 37th International Conference on Machine Learning (ICML)*, volume 119 of *PMLR*, 2020. URL <https://arxiv.org/abs/2002.04926>.
- [24] Javier García and Fernando Fernández. A comprehensive survey on safe reinforcement learning. *Journal of Machine Learning Research*, 16(42):1437–1480, 2015. URL <https://jmlr.org/papers/v16/garcia15a.html>.
- [25] Nan Geng, Tian Lan, Vaneet Aggarwal, Yuan Yang, and Mingwei Xu. A multi-agent reinforcement learning perspective on distributed traffic engineering. In *Proceedings of the 28th IEEE International Conference on Network Protocols (ICNP)*, pages 1–11, 2020. doi: 10.1109/ICNP49622.2020.9259413.
- [26] Fei Gui, Songtao Wang, Dan Li, Li Chen, Kaihui Gao, Congcong Min, and Yi Wang. RedTE: Mitigating subsecond traffic bursts with real-time and distributed traffic engineering. In *Proceedings of ACM SIGCOMM 2024*, pages 71–85, 2024. doi: 10.1145/3651890.3672231.
- [27] Peter Henderson, Riashat Islam, Philip Bachman, Joelle Pineau, Doina Precup, and David Meger. Deep reinforcement learning that matters. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 32, 2018. doi: 10.1609/aaai.v32i1.11694.
- [28] Chi-Yao Hong, Srikanth Kandula, Ratul Mahajan, Ming Zhang, Vijay Gill, Mohan Nanduri, and Roger Wattenhofer. Achieving high utilization with software-driven WAN. In *Proceedings of ACM SIGCOMM 2013*, pages 15–26, 2013. doi: 10.1145/2486001.2486012.
- [29] Oliver Hope and Eiko Yoneki. GDDR: GNN-based data-driven routing. In *Proceedings of the 41st IEEE International Conference on Distributed Computing Systems (ICDCS)*, pages 517–527, 2021. doi: 10.1109/ICDCS51616.2021.00056.
- [30] Shengyi Huang and Santiago Ontañón. A closer look at invalid action masking in policy gradient algorithms. In *Proceedings of the International FLAIRS Conference*, volume 35, 2022. doi: 10.32473/flairs.v35i.130584.
- [31] Sushant Jain, Alok Kumar, Subhasree Mandal, Joon Ong, Leon Poutievski, Arjun Singh, Subbaiah Venkata, Jim Wanderer, Junlan Zhou, Min Zhu, Jon Zolla, Urs Hölzle, Stephen Stuart, and Amin Vahdat. B4: Experience with a globally-deployed software defined WAN. In *Proceedings of ACM SIGCOMM 2013*, pages 3–14, 2013. doi: 10.1145/2486001.2486019.
- [32] Nan Jiang, Alex Kulesza, Satinder Singh, and Richard Lewis. The dependence of effective planning horizon on model accuracy. In *Proceedings of the 2015 International Conference on Autonomous Agents and Multiagent Systems (AAMAS)*, pages 1181–1189, 2015. URL <https://web.eecs.umich.edu/~baveja/Papers/gamma-AAMAS-final.pdf>. Best Paper Award.
- [33] Srikanth Kandula, Dina Katabi, Bruce Davie, and Anna Charny. Walking the tightrope: Responsive yet stable traffic engineering. In *Proceedings of ACM SIGCOMM 2005*, pages 253–264, 2005. doi: 10.1145/1080091.1080122.
- [34] Praveen Kumar, Yang Yuan, Chris Yu, Nate Foster, Robert Kleinberg, Petr Lapukhov, Chiun Lin Lim, and Robert Soulé. Semi-oblivious traffic engineering: The road not taken. In *Proceedings of the 15th USENIX Symposium on Networked Systems Design and Implementation (NSDI)*, 2018. URL <https://www.usenix.org/conference/nsdi18/presentation/kumar>.
- [35] Cassidy Laidlaw, Stuart Russell, and Anca Dragan. Bridging RL theory and practice with the effective horizon. In *Advances in Neural Information Processing Systems 36 (NeurIPS 2023)*, 2023. URL <https://arxiv.org/abs/2304.09853>. Oral. arXiv:2304.09853.
- [36] Cassidy Laidlaw, Banghua Zhu, Stuart Russell, and Anca Dragan. The effective horizon explains deep RL performance in stochastic environments. In *International Conference on Learning Representations (ICLR)*, 2024. URL <https://arxiv.org/abs/2312.08369>. Spotlight. arXiv:2312.08369.
- [37] John Langford and Tong Zhang. The epoch-greedy algorithm for multi-armed bandits with side information. In *Advances in Neural Information Processing Systems 20 (NIPS 2007)*, 2007. URL <https://proceedings.neurips.cc/paper/2007/hash/4b04a686b0ad13dce35fa99fa4161c65-Abstract.html>.
- [38] Ximeng Liu, Shizhen Zhao, Yong Cui, and Xinbing Wang. FIGRET: Fine-grained robustness-enhanced traffic engineering. In *Proceedings of ACM SIGCOMM 2024*, pages 117–135, 2024. doi: 10.1145/3651890.3672258.
- [39] Horia Mania, Aurelia Guy, and Benjamin Recht. Simple random search of static linear policies is competitive for reinforcement learning. In *Advances in Neural Information Processing Systems 31 (NeurIPS 2018)*, 2018. URL <https://papers.nips.cc/paper/2018/hash/7634ea65a4e6d9041cf3f7de18e334a-Abstract.html>.
- [40] Volodymyr Mnih, Koray Kavukcuoglu, David Silver, Andrei A. Rusu, Joel Veness, Marc G. Bellemare, Alex Graves, Martin Riedmiller, Andreas K. Fidjeland, Georg Ostrovski, Stig Petersen, Charles Beattie, Amir Sadik, Ioannis Antonoglou, Helen King, Dharshan Kumaran, Daan Wierstra, Shane Legg, and Demis Hassabis. Human-level control through deep reinforcement learning. *Nature*, 518(7540):529–533, 2015. doi: 10.1038/nature14236.
- [41] Andrew Y. Ng, Daishi Harada, and Stuart J. Russell. Policy invariance under reward transformations: Theory and application to reward shaping. In *Proceedings of the Sixteenth International Conference on Machine Learning (ICML)*, pages 278–287. Morgan Kaufmann, 1999.

- [42] Andrew Patterson, Samuel Neumann, Martha White, and Adam White. Empirical design in reinforcement learning. *Journal of Machine Learning Research*, 25(318):1–63, 2024. URL <https://www.jmlr.org/papers/v25/23-0183.html>.
- [43] Xinglong Pei, Penghao Sun, Yuxiang Hu, Dan Li, Bo Chen, and Le Tian. Enabling efficient routing for traffic engineering in SDN with deep reinforcement learning. *Computer Networks*, 241:110220, 2024. doi: 10.1016/j.comnet.2024.110220.
- [44] Yarin Perry, Felipe Vieira Frujeri, Chaim Hoch, Srikanth Kandula, Ishai Menache, Michael Schapira, and Aviv Tamar. DOTE: Rethinking (predictive) WAN traffic engineering. In *Proceedings of the 20th USENIX Symposium on Networked Systems Design and Implementation (NSDI)*, 2023. URL <https://www.usenix.org/conference/nsdi23/presentation/perry>. Best Paper Award. arXiv:2303.00735.
- [45] Marek Petrik and Bruno Scherrer. Biasing approximate dynamic programming with a lower discount factor. In *Advances in Neural Information Processing Systems 21 (NIPS 2008)*, pages 1265–1272, 2008. URL [https://proceedings.neurips.cc/paper\\_files/paper/2008/file/08c5433a60135c32e34f46a71175850c-Paper.pdf](https://proceedings.neurips.cc/paper_files/paper/2008/file/08c5433a60135c32e34f46a71175850c-Paper.pdf).
- [46] Antonin Raffin, Ashley Hill, Adam Gleave, Anssi Kanervisto, Maximilian Ernestus, and Noah Dormann. Stable-Baselines3: Reliable reinforcement learning implementations. *Journal of Machine Learning Research*, 22(268):1–8, 2021. URL <https://jmlr.org/papers/v22/20-1364.html>.
- [47] John Schulman, Philipp Moritz, Sergey Levine, Michael I. Jordan, and Pieter Abbeel. High-dimensional continuous control using generalized advantage estimation. In *International Conference on Learning Representations (ICLR)*, 2016. URL <https://arxiv.org/abs/1506.02438>.
- [48] John Schulman, Filip Wolski, Prafulla Dhariwal, Alec Radford, and Oleg Klimov. Proximal policy optimization algorithms, 2017. URL <https://arxiv.org/abs/1707.06347>.
- [49] Giorgio Stampa, Marta Arias, David Sanchez-Charles, Victor Muntès-Mulero, and Albert Cabellos. A deep-reinforcement learning approach for software-defined networking routing optimization, 2017. URL <https://arxiv.org/abs/1709.07080>.
- [50] Asaf Valadarsky, Michael Schapira, Dafna Shahaf, and Aviv Tamar. Learning to route. In *Proceedings of the 16th ACM Workshop on Hot Topics in Networks (HotNets-XVI)*, pages 185–191, 2017. doi: 10.1145/3152434.3152441. Extended version: arXiv:1708.03074, “A Machine Learning Approach to Routing”.
- [51] Hao Wang, Haiyong Xie, Lili Qiu, Yang Richard Yang, Yin Zhang, and Albert Greenberg. COPE: Traffic engineering in dynamic networks. In *Proceedings of ACM SIGCOMM 2006*, pages 99–110, 2006. doi: 10.1145/1159913.1159926.
- [52] Ning Wang, Kin Ho, George Pavlou, and Michael Howarth. An overview of routing optimization for Internet traffic engineering. *IEEE Communications Surveys & Tutorials*, 10(1):36–56, 2008. doi: 10.1109/COMST.2008.4483669.
- [53] Zhiying Xu, Francis Y. Yan, Rachee Singh, Justin T. Chiu, Alexander M. Rush, and Minlan Yu. Teal: Learning-accelerated optimization of WAN traffic engineering. In *Proceedings of ACM SIGCOMM 2023*, pages 378–393, 2023. doi: 10.1145/3603269.3604857.
- [54] Zhiyuan Xu, Jian Tang, Jingsong Meng, Weiye Zhang, Yanzhi Wang, Chi Harold Liu, and Dejun Yang. Experience-driven networking: A deep reinforcement learning based approach. In *Proceedings of IEEE INFOCOM 2018*, 2018. doi: 10.1109/INFOCOM.2018.8485853.
- [55] Francis Y. Yan, Hudson Ayers, Chenzhi Zhu, Sadjad Fouladi, James Hong, Keyi Zhang, Philip Levis, and Keith Winstein. Learning in situ: a randomized experiment in video streaming. In *Proceedings of the 17th USENIX Symposium on Networked Systems Design and Implementation (NSDI)*, pages 495–511, 2020. URL <https://www.usenix.org/conference/nsdi20/presentation/yan>.
- [56] Junjie Zhang, Minghao Ye, Zehua Guo, Chen-Yu Yen, and H. Jonathan Chao. CFR-RL: Traffic engineering with reinforcement learning in SDN. *IEEE Journal on Selected Areas in Communications*, 38(10):2249–2259, 2020. doi: 10.1109/JSAC.2020.3000371.
- [57] Zihan Zhang, Xiangyang Ji, and Simon S. Du. Is reinforcement learning more difficult than bandits? A near-optimal algorithm escaping the curse of horizon. In *Proceedings of the 34th Conference on Learning Theory (COLT)*, volume 134 of *PMLR*, pages 4528–4531, 2021. URL <https://proceedings.mlr.press/v134/zhang21b.html>.

## A Environment details

All values are generated from the configuration files by `scripts/study/analyze.py`. The topology has 18 routers, 32 bidirectional links (64 directed), 45.0 Gb/s of total directed capacity, and 17 demands whose base (profile-peak) rates sum to 2.49 Gb/s.

**Table 9:** Traffic classes. Priority weights  $q_d = \pi_d/6$  enter the SLA and disconnection terms; protected classes may only be moved onto candidates whose projected gross bottleneck utilization is at most 100 %.

| Class      | Priority | Delay SLA (ms) | Loss SLA (%) | Protected | Profile       | Demands |
|------------|----------|----------------|--------------|-----------|---------------|---------|
| voice      | 5        | 60             | 0.5          | yes       | business      | 3       |
| video      | 4        | 120            | 1            | no        | evening_video | 4       |
| vpn        | 3        | 150            | 1            | no        | business      | 3       |
| besteffort | 2        | 250            | 3            | no        | consumer      | 3       |
| bulk       | 1        | 400            | 5            | no        | night_bulk    | 3       |
| critical   | 6        | 80             | 0.2          | yes       | flat          | 1       |

**Table 10:** Scenarios. Learners train only on `random_day`, whose bursts, single link failure (probability 0.5) and flash crowd (probability 0.3) are drawn per episode; all seven evaluation scenarios are scripted and differ from the training distribution.

| Scenario                | Start (h) | Decisions | Multiplier | Noise $\sigma$ | Events             |
|-------------------------|-----------|-----------|------------|----------------|--------------------|
| full_day                | 0         | 288       | 1          | 1              | none               |
| evening_peak            | 16        | 84        | 1.15       | 1              | none               |
| flash_crowd             | 18        | 60        | 1.1        | 1              | flash crowd        |
| link_failure            | 17        | 60        | 1.1        | 1              | link down, link up |
| deceptive_local_optimum | 9         | 60        | 1          | 0.8            | burst              |
| ood_double_failure      | 10        | 60        | 1.1        | 1              | link down, link up |
| overload_stress         | 18        | 48        | 1.6        | 1              | none               |
| random_day              | 0         | 288       | 1          | 1              | randomized         |

**Observation (604 features, feature-major, all in  $[0, 1]$ ).** Per directed link:  $\text{sat}(\rho_e) = \rho_e / (1 + \rho_e)$  and the up flag ( $2 \times 64$ ). Per demand ( $28 \times 17$ ):  $\text{sat}(\text{offered}/\text{base})$ , priority/6, protected flag,  $\text{sat}(\text{delay}/\text{SLA})$ ,  $\text{sat}(\text{loss}/\text{SLA})$ , path age/12 (capped), dwell/3, disconnected flag, one-hot current path (4), one-hot previous TE path (4), candidate liveness (4), candidate propagation delay/100 ms (4), and  $\text{sat}$  of the projected gross bottleneck utilization of moving to each candidate (4). Time of day, episode progress and future events are deliberately excluded.

**Episode boundaries.** Events scheduled at minute  $t$  are applied in the right-closed window  $(t - 1, t]$ , so the observation at a boundary reflects every event up to that time. Episodes never terminate; they are truncated at the scenario end, so paired controllers have identical horizons.

## B Hyper-parameters and budgets

**Table 11:** Learner configurations. All values are read from `configs/training.yaml` (PPO) and `configs/experiments/learning_v2.yaml` (bandit/Q); the resolved configuration of every run is stored in its manifest.

|                       | MaskablePPO                                             | Masked bandit / Q-learner                                             |
|-----------------------|---------------------------------------------------------|-----------------------------------------------------------------------|
| Network               | separate [256, 256] tanh (policy, value)                | 604–256–256–69 ReLU                                                   |
| Optimizer             | Adam, $3 \cdot 10^{-4}$ , grad-norm clip 0.5            | Adam, $3 \cdot 10^{-4}$ , grad-norm clip 1.0                          |
| Discount              | $\gamma = 0.995$ , GAE $\lambda = 0.95$                 | $\gamma = 0$ (bandit); $\{0.5, 0.9, 0.99\}$ (Q)                       |
| Data per update       | rollout $16 \times 512 = 8192$ , 8 epochs, batch 512    | replay $10^5$ , batch 512, every 64 transitions                       |
| Gradient steps / 400k | 6,144                                                   | 6,187                                                                 |
| Exploration           | policy entropy, coef. 0.01                              | masked $\epsilon$ -greedy, $0.2 \rightarrow 0.02$ over $2 \cdot 10^5$ |
| Loss                  | clipped surrogate $(0.2) + 0.5 \times \text{value MSE}$ | Huber on selected head                                                |
| Target                | GAE returns                                             | $r$ (bandit); Eq. (4), target copy every 250 updates (Q)              |
| Masking               | logits $\rightarrow -10^8$                              | scores $\rightarrow -\infty$ ; random actions from legal set          |
| Warm-up               | –                                                       | 4,096 transitions                                                     |

**PPO sensitivity grid (E3).** One factor at a time around the default, on training root 42, selected on validation seeds only: learning rate  $\{10^{-4}, 10^{-3}\}$ ; rollout length per environment 2048 (batch 1024); entropy coefficient  $\{0, 0.03\}$ ; clip range 0.1; and, added after the return decomposition because they target the proposed mechanism (noisy advantages hiding small immediate costs) while keeping  $\gamma = 0.995$ , reward normalization (running return scale, observations untouched) and GAE  $\lambda = 0.8$ . Two planned variations (rollout 128, network [64, 64]) were dropped for compute before they ran. The discount is varied separately (E2,  $\gamma \in \{0, 0.9\}$ ). The configuration with the best validation return is retrained on the four other primary roots (E3b) and compared with the bandit and the default on the same roots.

**Table 12:** PPO sensitivity on root 42 (E3). Validation: mean return of the best checkpoint on validation seeds (the selection criterion). Test: the selected checkpoint on the 140 test episodes, reported for completeness only. “ppo” is the default configuration, “bandit” the reference.

| Configuration | Validation | Test  |
|---------------|------------|-------|
| ent003        | 15.3       | 20.7  |
| gae08         | 13.7       | 19.9  |
| bandit        | 12.2       | 18.3  |
| rewnorm       | −5.1       | 2.5   |
| ppo           | −19.1      | −13.7 |
| lr1e4         | −23.9      | −18.6 |
| lr1e3         | −24.7      | −17.6 |
| ent0          | −26.2      | −20.4 |
| clip01        | −27.8      | −22.2 |
| ns2048        | −28.2      | −22.6 |

## C Additional results

**Table 13:** Primary comparison per training root (mean test return over 140 episodes): selected checkpoint (validation seeds) and final 400k checkpoint.

| Root   | Masked bandit |        |       | MaskablePPO |        |       |
|--------|---------------|--------|-------|-------------|--------|-------|
|        | Selected      | Return | Final | Selected    | Return | Final |
| 42     | 350k          | 18.3   | 20.3  | 350k        | −13.7  | −47.4 |
| 141421 | 350k          | 25.0   | 18.6  | 300k        | 14.8   | −9.4  |
| 161803 | 400k          | 28.5   | 28.5  | 300k        | 5.8    | −4.6  |
| 271828 | 350k          | 28.4   | 24.1  | 400k        | −22.5  | −22.5 |
| 314159 | 250k          | 25.8   | 16.4  | 200k        | 16.9   | 13.2  |

## D Preregistered hypotheses and outcomes

Hypotheses H0–H8 were committed before any new learner result existed (`docs/HYPOTHESES.md`); H9–H14 were added later, each before its experiment ran, and are exploratory. Three of the five primary training roots are the reproduction runs, whose direction was known from the closed study when H1 was written; the two additional roots were new. The frozen-exogenous rollout was added after an early look at the clairvoyant diagnostic.

## E Reproducibility

**Software and hardware.** Python 3.13, PyTorch 2.14 (CPU build), Stable-Baselines3 and SB3-contrib 2.9, SciPy HiGHS for MILP-track; exact versions are pinned in `requirements/lock-cpu-py313.txt`. All post-V2 results were produced on one 4-core x86-64 CPU machine without a GPU, with one PyTorch thread per training process and at most four concurrent training processes. Training a learner for 400k transitions took 0.5–1 h of wall time on this machine.

**Determinism.** Episode seeds are derived from the training root (training) or listed explicitly (evaluation); traffic is generated from per-episode generators that no routing decision consumes, so every controller sees identical traffic on a given episode. On CPU with one thread a training run is deterministic given its root; the study trainer reproduces the closed study’s trainer bit for bit (a slow test checks the network parameters). The closed study’s own learner runs used CUDA and are therefore reproduced statistically, not exactly (§6.1); its baselines reproduce exactly.

**Table 14:** Learner minus reference, paired test returns (two-stage bootstrap 95 % CI over roots and episodes) and the number of roots on which the difference is positive.

| Learner       | Reference             | Difference [95 % CI] | Roots > 0 |
|---------------|-----------------------|----------------------|-----------|
| Masked bandit | No-op                 | 120.7 [117.0, 123.6] | 5/5       |
| Masked bandit | Static SP             | 121.7 [118.0, 124.6] | 5/5       |
| Masked bandit | Random valid          | 132.7 [128.7, 136.1] | 5/5       |
| Masked bandit | CSPF                  | 50.8 [47.1, 53.7]    | 5/5       |
| Masked bandit | Greedy                | 21.2 [17.6, 24.1]    | 5/5       |
| Masked bandit | MILP-track            | 1.2 [-2.5, 4.0]      | 4/5       |
| Masked bandit | Oracle-1 <sup>†</sup> | -21.7 [-25.4, -18.9] | 0/5       |
| MaskablePPO   | No-op                 | 95.7 [81.9, 109.2]   | 5/5       |
| MaskablePPO   | Static SP             | 96.7 [82.9, 110.2]   | 5/5       |
| MaskablePPO   | Random valid          | 107.8 [93.8, 121.2]  | 5/5       |
| MaskablePPO   | CSPF                  | 25.9 [12.0, 39.4]    | 5/5       |
| MaskablePPO   | Greedy                | -3.7 [-17.6, 9.8]    | 3/5       |
| MaskablePPO   | MILP-track            | -23.8 [-37.6, -10.3] | 0/5       |
| MaskablePPO   | Oracle-1 <sup>†</sup> | -46.6 [-60.3, -33.2] | 0/5       |

**Table 15:** Masked bandit – MILP-track (untuned default) per scenario on the test episodes: two-stage bootstrap 95 % CI and number of roots with a positive difference.

| Scenario       | Difference [95 % CI] | Roots > 0 |
|----------------|----------------------|-----------|
| Deceptive      | 4.6 [1.9, 7.0]       | 4/5       |
| Evening peak   | 8.9 [6.5, 11.3]      | 5/5       |
| Flash crowd    | 2.6 [0.3, 4.9]       | 4/5       |
| Full day       | 9.8 [2.6, 15.9]      | 4/5       |
| Link failure   | 1.6 [-7.1, 9.5]      | 3/5       |
| Double failure | -16.6 [-23.9, -9.5]  | 0/5       |
| Overload       | -2.8 [-5.8, 0.2]     | 1/5       |

**One-command entry points.** `python scripts/reproduce.py -suite sanity` (tests, exact baseline reproduction, a short training and evaluation of both learners; minutes), `-suite core` (one further 400k root of both learners and all reference policies; hours) and `-suite full` (every registered run; about a day on the machine above). `-suite analysis` rebuilds every table, every number in this paper (`paper/generated/numbers.tex`), the figures and the PDF from the stored per-episode results.

**Provenance.** Every experiment is declared in `experiments/registry/*.yaml` before it runs; each run writes a manifest with the git commit, resolved configuration, environment specification, seeds, integrity counters and wall time. Per-episode test results of every checkpoint that was evaluated, the mask audit and the sequentiality diagnostic are versioned under `experiments/raw/`. Deviations from the preregistered plan (`docs/HYPOTHESES.md`) and failed or interrupted runs are recorded in `docs/EXPERIMENT_LOG.md`; interrupted runs were restarted from scratch with the same root and their partial outputs kept.

**Table 16:** Sequentiality diagnostic per scenario at  $H = 24$  (states with at least 24 remaining intervals): agreement of the myopic and the 24-interval best action, share of the 24-interval gain captured by the myopic action, and share of states in which the 24-interval best action has negative immediate advantage, under clairvoyant and frozen-exogenous rollouts (%).

| Scenario       | States | Clairvoyant |          |           | Frozen exogenous |          |           |
|----------------|--------|-------------|----------|-----------|------------------|----------|-----------|
|                |        | Agree       | Captured | Sacrifice | Agree            | Captured | Sacrifice |
| Full day       | 201    | 36          | 49       | 60        | 91               | 96       | 7         |
| Evening peak   | 48     | 29          | 42       | 56        | 85               | 97       | 12        |
| Flash crowd    | 30     | 23          | 65       | 60        | 83               | 99       | 7         |
| Link failure   | 30     | 30          | 57       | 57        | 93               | 99       | 3         |
| Deceptive      | 30     | 40          | 52       | 53        | 87               | 92       | 13        |
| Double failure | 30     | 23          | 36       | 73        | 93               | 97       | 7         |
| Overload       | 21     | 24          | 67       | 38        | 67               | 96       | 5         |

**Table 17:** Inference cost per decision (process CPU time, one thread, 60 states of a test episode) and training cost (400k transitions, wall time and throughput measured with 3–6 concurrent jobs on 4 cores; the simulator dominates).

| Policy                | Parameters | Train (min) | Transitions/s | ms / decision | p95 ms  |
|-----------------------|------------|-------------|---------------|---------------|---------|
| Masked bandit         | 238,405    | 45          | 147           | 0.112         | 0.156   |
| MaskablePPO           | 459,334    | 49          | 136           | 0.335         | 0.362   |
| MILP-track            | —          | —           | —             | 31.7          | 65.7    |
| Greedy                | —          | —           | —             | 0.0639        | 0.0758  |
| CSPF                  | —          | —           | —             | 0.0412        | 0.254   |
| Static SP             | —          | —           | —             | 0.00342       | 0.00412 |
| Oracle-1 <sup>†</sup> | —          | —           | —             | 359           | 590     |

**Table 18:** Outcome of every hypothesis. “Pending”: the experiment had not finished when this version was compiled.

| ID  | Hypothesis (abridged)                                                                              | Outcome                                                                                                                                                                                                                      |
|-----|----------------------------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| H0  | Closed-study bandit > PPO reproduces                                                               | Direction reproduced on 3/3 roots; magnitude not (30.1 vs. 9.2)                                                                                                                                                              |
| H1  | Gap survives new protocol and more roots                                                           | Supported: 24.9, $t$ -interval [3.4, 46.4], 5/5 roots                                                                                                                                                                        |
| H2  | Shortening PPO’s horizon does not hurt it                                                          | Supported: PPO $\gamma = 0$ 23.6 vs. $\gamma = 0.995$ −6.4                                                                                                                                                                   |
| H3  | Bootstrapping does not help a value learner                                                        | Supported ( $\gamma = 0.5$ : 1.4; $\gamma = 0.9$ : −4.8, 0/3 roots better; $\gamma = 0.99$ : −20.9 (0/3 roots better) )                                                                                                      |
| H4  | PPO deficit is not a tuning artefact                                                               | Root 42 (selection root): best configuration ent003 validation 15.3 vs. default −19.1 and bandit 12.2; test 20.7 vs. bandit 18.3 (falsified on this root). Fresh roots (E3b, 4): tuned PPO − bandit = −28.6, positive on 0/4 |
| H5  | Non-myopic structure exists, mostly from (a) one-off cost vs. persistent gain and (b) clairvoyance | Prediction met (frozen agreement 88 % > live 33 %; Oracle-3/6 gains small). Component (a) <i>not</i> supported: cost amortization changes the best move in only 12 % of states                                               |
| H6  | Under delay the bandit fails and sequential learners do better                                     | Bandit fails (22.0 → −30.3); PPO better on 2/2 roots; Q $\gamma = 0.9$ −5.6 (2/2 roots better). No learner reaches MILP-track (11.3)                                                                                         |
| H7  | PPO does not catch up with 3× budget                                                               | pending                                                                                                                                                                                                                      |
| H8  | The bandit is closer to the myopic optimum than PPO; PPO does not find non-myopic value            | Supported, though neither learner is close to the optimum (Table 8)                                                                                                                                                          |
| H9  | Time of day helps Q $\gamma = 0.9$ more than the bandit                                            | Not supported: clock adds +0.4 to Q $\gamma = 0.9$ and +1.9 to the bandit (2 roots)                                                                                                                                          |
| H10 | PPO $\gamma = 0$ removes most flapping                                                             | Supported: reversals 52.8 → 1.0 per episode                                                                                                                                                                                  |
| H11 | The diagnostic flags the delayed regime (frozen agreement drops)                                   | pending (19 of 21 episodes)                                                                                                                                                                                                  |
| H12 | A reactive continuation raises open-loop agreement                                                 | pending (0 of 21 episodes)                                                                                                                                                                                                   |